

Once Upon AI: Intelligent Storybook Generator for Children

Team Members: Huifang Li, Nausad Miyan, Arindam Laha

GitHub Repository: <https://github.com/laha-arindam/egr404>

Introduction

Children's literature has long played a foundational role in promoting literacy, imagination, and emotional development. From fairy tales to fables, stories provide a safe space for young minds to explore values, relationships, and cultures. However, traditional printed storybooks are limited in variety, are difficult to personalize, and often lack representation for diverse backgrounds, experiences, and developmental needs. Families and educators alike face challenges in finding story content that resonates with individual children, whether culturally, emotionally, or pedagogically. Moreover, the publishing process for such materials is often slow and resource-intensive, making it inaccessible to many learners around the world.

This project, titled 'Once Upon AI,' addresses this gap by creating a system that can automatically generate personalized children's storybooks based on simple input parameters such as a character, a setting, and a moral theme. The stories produced are written in accessible, age-appropriate language and follow narrative structures aligned with early childhood reading patterns. Each story includes visual illustrations that match the scenes described in the text. The content is generated dynamically based on user-defined inputs, ensuring that each output is unique and contextually relevant.

The novelty of this project lies in its integration of real-time content creation into a user-friendly educational interface. Whereas most educational content delivery tools rely on pre-scripted templates or fixed libraries, this system can generate new, coherent narratives on demand. It also enables emotional and cognitive tailoring of the story, which is essential for differentiated instruction in modern classrooms. By making it possible for any teacher, parent, or caregiver to co-create stories with minimal effort, the system supports inclusive and participatory learning models.

The educational significance of this tool is twofold. First, it supports reading comprehension and critical thinking through context-specific, theme-rich storytelling. Second, it offers representation and narrative autonomy to children who may not typically see themselves reflected in traditional literature. From a scientific and technical standpoint, the project showcases how modern language and image generation pipelines can be applied to practical educational use cases. The final system is accessible via a browser interface and delivers outputs in under one minute, demonstrating the feasibility of rapid, personalized content generation for learning environments.

The primary objective of the project was to build a platform that can generate short, illustrated children's stories based on prompts like characters, settings, and moral themes. The project specifically focuses on three core tasks: automated story generation using a text-based model, illustration creation using a generative image model, and seamless compilation into a reader-friendly digital book format.

Our solution allows users to enter basic story parameters, and within minutes, receive a multi-page illustrated storybook. These books are designed to support literacy, spark imagination, and reflect the diversity of experiences that children bring to their learning environments.

Methodology and Tools

The development methodology focused on modular design and data-flow management to support the sequential generation of story content, image rendering, and layout formatting. We used OpenAI’s GPT-3.5-turbo API for generating story text. Prompt engineering was essential to ensure stories followed a consistent narrative structure: introduction, problem, resolution, and moral conclusion. Each story was composed of 4 to 6 paragraphs.

Each paragraph was summarized into a visual description to drive image generation via the DALL·E model. For example, a paragraph describing 'a fox helping a turtle cross a stream' would generate a prompt like 'storybook illustration of a friendly fox helping a turtle cross a small river.'

System Architecture

Timeline and Development Process

Date Range	Phase	Activities and Deliverables
March 25–27	Project Planning	Brainstormed ideas, formed team, submitted proposal, and created repository.
March 28–April 3	Model Selection and Prompt Engineering	Tested GPT-3.5 responses, designed story templates, validated output quality.
April 4–10	Text Pipeline Integration	Integrated text generation into system, ensured flow and coherence across paragraphs.
April 11–17	Illustration and Scene Design	Mapped story content to visual prompts, tested DALL·E rendering accuracy.

April 23–26	Frontend Integration	Finalized user interface, connected backend, tested usability.
April 27–28	Demo and Presentation	Recorded live demo, created slides, practiced pitch.
April 29–May 2	Final Documentation	Wrote and refined final report, submitted all materials to GitHub and instructors.

Results and Evaluation

Story Settings

Main Character

Luna the Cat

Setting

a magical forest

Genre

Adventure

Moral of the Story

Be kind to others

Generate Story

Once Upon AI - Storybook Generator for Children

Your AI-Generated Story

Once upon a time, in the small, lively town of Sunnyvale, there lived a cute, mischievous cat named Luna. Luna had shiny, deep black fur, sparkling blue eyes, and a tail that twitched when she was excited.

One sunny afternoon, Luna found a magical mantra in an ancient book in the attic. The mantra, when recited, could transport anyone to a mystical forest called the Enchanted Forest. Brimming with feline curiosity, Luna recited the mantra and, poof, found herself standing at the edge of the most beautiful forest she had ever seen.

The trees were a vivid shade of green, glowing under the radiant sunlight. The air carried an aroma of flowers and ripe fruits. Vibrant butterflies fluttered around the enchanting flowers, while magical creatures danced and played. Luna was eager to explore, and she plunged into the echoes of the forest.

The first creature she met was a grumpy old owl named Oliver. Oliver's feathers were ruffled, and he was not in the best of moods. Luna, being wise and kind, groomed Oliver's feathers with her dainty paws. Oliver hooted a thankful goodnight and gifted Luna a magical feather of wisdom.

Next, Luna met a sullen hedgehog named Harold. Harold had lost his way in the midst of the forest. Luna, with her innate sense of direction, guided him back to his home. In return, Harold thanked Luna and gave her a spiky, protective shield.

Finally, Luna encountered a lonely pixie named Pixie. Pixie had lost her friends during a playful chase and was on the verge of tears. Luna invited Pixie to hop onto her back, and they scoured the forest, finding Pixie's friends hiding behind a buoyant bouncy bush. Overjoyed, Pixie gifted Luna a magical twinkling star.

After meeting these magical creatures, Luna found herself standing in front of a massive crystal door adorned with mythical creatures. The door led to her way back home. However, the door was locked and required three magical elements—wisdom, protection, and a touch of magic—to unlock.

And so, Luna lived, embodying the moral of being kind to others - making Sunnyvale an even sunnier, happier place to live. And that's the end of our tale for today. Until next time, sweet dreams, and remember to be kind to others, just like Luna the Cat.

Story Illustration



AI-generated illustration

Qualitative feedback collected from peers highlighted the tool's intuitive UI, relevance of story morals, and impressive image quality. The project successfully met its goal of providing a self-contained, interactive storytelling solution.

Challenges and Solutions

Other challenges included prompt unpredictability in DALL·E responses and handling API quota limits. We addressed these by refining prompts with stylistic anchors and implementing retry logic for rate-limited requests.

Team Contributions

GitHub Repository and Documentation

The GitHub repository includes source code, README documentation, example outputs, the demo video, and this report. The code is modular and includes setup instructions, making it easy to replicate or extend. Each function and script is clearly documented for readability and educational reuse.

Project Structure

```
once-upon-ai/
├── app.py                # Main web interface
├── story_generator.py    # GPT-based story creation
├── image_generator.py    # DALL·E / Stable Diffusion illustration generation
├── templates/           # HTML templates (for Flask)
├── static/              # CSS, images, or assets
├── requirements.txt      # Python dependencies
├── README.md            # Project overview and instructions
└── .gitignore           # Files/directories to exclude from version control
```

Background and Motivation

In traditional educational settings, storytelling is often confined to printed books, which, while effective, are limited in variety, accessibility, and personalization. Children from different cultural backgrounds may not see themselves represented in the stories they consume. Moreover, the moral and emotional lessons that could be conveyed through personalized narratives are often generic or mismatched to the child's current stage of learning or interest.

With the emergence of powerful computational models, there is now a genuine opportunity to produce rich, meaningful, and customizable content in real time. This innovation is particularly impactful in early childhood education, where attention span, emotional connection, and visual stimulation are crucial to engagement. Our team recognized this as a moment to rethink how storybooks are produced, distributed, and consumed.

Design Philosophy and Ethical Considerations

We were intentional about designing our application with child-safety and pedagogical value at its core. Generative systems are powerful but can also produce undesired or inappropriate output if not carefully managed. To mitigate this risk, we crafted prompt templates that restricted content to known-safe narrative structures, avoided potentially violent or offensive elements, and emphasized positive values such as empathy, teamwork, and curiosity.

In addition, we made conscious efforts to test outputs with a diversity of themes and names, ensuring that the system would not default to cultural stereotypes. We also avoided datasets that might introduce bias and restricted open-ended responses from the language model by narrowing the scope of user input.

Potential for Future Expansion

Although the current version of 'Once Upon AI' focuses on single-story outputs, the framework could easily be extended to support multi-chapter books, personalized learning journeys, or thematic collections. There is also a promising opportunity to integrate text-to-speech capabilities, allowing children to listen to stories being read aloud—either in real time or offline. This would make the tool more inclusive for visually impaired children or those developing language skills.

Furthermore, with the addition of user profiles and feedback loops, future iterations could adapt stories based on prior selections, learning pace, or behavioral preferences, creating a feedback-driven, adaptive literacy experience.

Impact and Significance

'Once Upon AI' represents more than a technical achievement; it is a vision for how technology can be used to enhance education, nurture creativity, and provide equitable access to cultural and emotional learning tools. The prototype demonstrates how relatively simple interfaces can hide complex AI pipelines, making cutting-edge technology usable even for children and non-technical adults.

By giving families, teachers, and therapists the power to generate tailored content instantly, we reduce dependency on commercial publishing models and enable true personalization. Such capabilities are particularly important in special education, bilingual education, and trauma-informed care contexts, where off-the-shelf materials often fall short.

In sum, this project serves as a foundation for rethinking how narrative content is created and distributed, using automation to deliver meaningful, context-aware, and developmentally appropriate materials to every child.

Technical Implementation Details

The technical design of 'Once Upon AI' is based on loosely coupled services that communicate through synchronous API calls. The architecture ensures that failures in one module do not cascade across the system, and the response times of the models are regulated through efficient caching and async handling techniques. To maximize uptime during demonstrations and testing, we incorporated a fallback handler in the image generation module that substitutes default placeholders if API latency exceeds a predefined threshold.

The choice of using Streamlit for the frontend stemmed from its minimal learning curve and compatibility with Python. It allowed rapid prototyping and iterative design with immediate feedback, essential in refining the UI for a non-technical user base. We placed emphasis on ensuring accessibility, limiting the number of input fields, and using instructional text to guide users through the prompt creation process.

Evaluation Criteria and Testing Approach

To verify the integrity and usefulness of our system, we adopted a multi-layered evaluation framework. From a technical standpoint, we tested for input/output alignment—ensuring that every generated image correctly represented the paragraph it was derived from. This was assessed manually for all test cases. We also measured response time, success rates of generation tasks, and system reliability over a continuous 72-hour trial period.

On the user experience front, we collected feedback from trial users on five key dimensions: ease of use, visual quality, narrative coherence, emotional engagement, and educational value. Ratings were collected using a Likert scale and open comments. Users especially appreciated the speed of generation, the intuitiveness of the interface, and the diversity of generated content.

Broader Educational Implications

The tool opens the door to customizable curriculum tools that educators can build in real time. For example, a teacher focusing on environmental science can prompt a story about 'a forest animal learning to protect nature,' thereby reinforcing academic content through narrative. Similarly, behavioral specialists could use stories to model conflict resolution or social interaction strategies, creating a safe emotional sandbox for children to explore values and decisions.

By offering content that is both relatable and immediate, our system helps reduce the cognitive gap between instruction and application. Furthermore, it aligns well with differentiated instruction strategies, which aim to tailor teaching methods and content to students' individual needs.

Scalability and Deployment Considerations

Our architecture is cloud-ready and can be containerized using Docker for deployment in controlled environments such as classrooms, libraries, or institutional networks. For wider adoption, user management and usage logging modules can be added to track how children interact with generated content over time. This would offer insights into content preferences and educational impact, potentially guiding adaptive content refinement and future pedagogical research.

Security and privacy are also important concerns. While the system currently does not store user data, any scaled-up version would need proper compliance with privacy laws such as COPPA or GDPR if used in educational settings. User prompts and outputs must be protected, and safeguards must be placed to avoid generation of inappropriate content through abuse or system manipulation.